

## Code for My heart will go on

```
[1]: import numpy as np
import numpy.typing as npt
import pandas as pd
```

```
[2]: train_url = "http://s3.amazonaws.com/assets.datacamp.com/course/Kaggle/
→train.csv"
train = pd.read_csv(train_url) #training set
test_url = "http://s3.amazonaws.com/assets.datacamp.com/course/Kaggle/test.
→csv"
test = pd.read_csv(test_url) #test set
```

```
[3]: train.describe()
```

```
[3]:
```

	PassengerId	Survived	Pclass	Age	SibSp	\
count	891.000000	891.000000	891.000000	714.000000	891.000000	
mean	446.000000	0.383838	2.308642	29.699118	0.523008	
std	257.353842	0.486592	0.836071	14.526497	1.102743	
min	1.000000	0.000000	1.000000	0.420000	0.000000	
25%	223.500000	0.000000	2.000000	20.125000	0.000000	
50%	446.000000	0.000000	3.000000	28.000000	0.000000	
75%	668.500000	1.000000	3.000000	38.000000	1.000000	
max	891.000000	1.000000	3.000000	80.000000	8.000000	

  

	Parch	Fare
count	891.000000	891.000000
mean	0.381594	32.204208
std	0.806057	49.693429
min	0.000000	0.000000
25%	0.000000	7.910400
50%	0.000000	14.454200
75%	0.000000	31.000000
max	6.000000	512.329200

### T8

```
[4]: print('median of age is', age_med := train['Age'].median())
```

median of age is 28.0

```
[5]: train['Age'] = train['Age'].fillna(age_med)
```

### T9

```
[6]: print('Embarked Mode is', embark_mode := train['Embarked'].mode()[0])
```

Embarked Mode is S

```
[7]: train['Embarked'] = train['Embarked'].fillna(embark_mode)
train.loc[train["Embarked"] == "S", "Embarked"] = 0
train.loc[train["Embarked"] == "C", "Embarked"] = 1
train.loc[train["Embarked"] == "Q", "Embarked"] = 2
```

```
[8]: print('Sex Mode is', sex_mode := train['Sex'].mode()[0])
```

Sex Mode is male

```
[9]: train['Sex'] = train['Sex'].fillna(sex_mode)
      train.loc[train["Sex"] == "male", "Sex"] = 0
      train.loc[train["Sex"] == "female", "Sex"] = 1
```

## T10, T11

```
[10]: class LogisticRegressionGradient:
      def __init__(self, lr=0.00001, random_state=42, epochs=10_000, threshold=0.
      ↪5):
          self.lr = lr
          self.random_state = random_state
          self.epochs = epochs
          self.threshold = threshold

          @staticmethod
          def logist(X: np.array):
              X = np.clip(X, -600, 600) # for overflow
              mask = X >= 0
              X[mask] = np.exp(X[mask]) / (1 + np.exp(X[mask]))
              X[~mask] = 1 / (1 + np.exp(-X[~mask]))
              return X

          def fit(self, X: npt.ArrayLike, y: npt.ArrayLike):
              X = np.array(X)
              y = np.array(y)

              np.random.seed(self.random_state)
              X = np.hstack((np.ones(X.shape[0]).reshape(-1, 1), X)) # add bias

              self.params = np.random.randn(X.shape[1])

              for _ in range(self.epochs):
                  y_pred = self.logist(X @ self.params)
                  diff = y - y_pred
                  loss = X.T @ diff

                  self.params += self.lr * loss

              return self

          def predict(self, X: npt.ArrayLike):
              X = np.array(X)
              X = np.hstack((np.ones(X.shape[0]).reshape(-1, 1), X)) # add bias
              return (self.logist(X @ self.params) >= self.threshold).astype(int)
```

```
[11]: X = np.array(train[["Pclass", "Sex", "Age", "Embarked"]].values, dtype = np.
      ↪float64)
      y = np.array(train['Survived'], dtype=np.float64)
```

```
[12]: lr = LogisticRegressionGradient()
      lr.fit(X, y)
```

```
[12]: <__main__.LogisticRegressionGradient at 0x1287b3c40>
```

```
[13]: test['Age'] = test['Age'].fillna(age_med)

      test['Embarked'] = test['Embarked'].fillna(embark_mode)
```

```
test.loc[test["Embarked"] == "S", "Embarked"] = 0
test.loc[test["Embarked"] == "C", "Embarked"] = 1
test.loc[test["Embarked"] == "Q", "Embarked"] = 2

test['Sex'] = test['Sex'].fillna(sex_mode)
test.loc[test["Sex"] == "male", "Sex"] = 0
test.loc[test["Sex"] == "female", "Sex"] = 1
```

```
[14]: y_pred = lr.predict(np.array(test[["Pclass", "Sex", "Age", "Embarked"]],
↳dtype=float))

pd.DataFrame({
    'PassengerId': test['PassengerId'],
    'Survived': y_pred
}).to_csv('Submit.csv', index=False)
```

## T12

```
[15]: def accuracy_score(y_test, y_pred):
    if y_test.shape[0] != y_pred.shape[0]:
        raise ValueError("Shape are not equal")
    return (y_test == y_pred).sum() / y_test.shape[0]
```

```
[16]: y_pred = lr.predict(X)
print('Accuracy score of training set is', accuracy_score(y, y_pred))
```

Accuracy score of training set is 0.7968574635241302

Add high order feature ( $x_1, x_1^2, x_2 \dots$ )

```
[17]: train['Age_squared'] = train['Age'] ** 2
test['Age_squared'] = test['Age'] ** 2
train['Age_Cubic'] = train['Age'] ** 3
test['Age_Cubic'] = test['Age'] ** 3

X_ho_train = np.array(train[["Pclass", "Sex", "Age", "Age_squared",
↳"Age_Cubic", "Embarked"]].values, dtype = np.float64)
X_ho_test = np.array(test[["Pclass", "Sex", "Age", "Age_squared",
↳"Age_Cubic", "Embarked"]].values, dtype = np.float64)

lr_ho = LogisticRegressionGradient().fit(X_ho_train, y)
y_pred_ho_train = lr_ho.predict(X_ho_train)

print(lr_ho.params)

print('Accuracy score of training set with high order feature is',
↳accuracy_score(y, y_pred_ho_train))
```

```
[ 0.70055359 -12.32966498 11.8377913    9.44359665 318.83403664
 -87.10590108  4.65074534]
```

Accuracy score of training set with high order feature is 0.6273849607182941

```
[18]: y_pred_ho_test = lr_ho.predict(X_ho_test)
pd.DataFrame({
    'PassengerId': test['PassengerId'],
    'Survived': y_pred_ho_test
}).to_csv('Submit_highorder.csv', index=False)
```

## T13

```
[19]: X_train = np.array(train[["Sex", "Age"]].values, dtype = np.float64)
      X_test = np.array(test[["Sex", "Age"]].values, dtype = np.float64)

      lr_sa = LogisticRegressionGradient().fit(X_train, y)
      y_pred_sa_train = lr_sa.predict(X_train)

      print(lr_sa.params)
      print('Accuracy score of training set with only Sex and Age is',
      ↪accuracy_score(y, y_pred_sa_train))

      [-1.01863706  2.34645073 -0.01149691]
      Accuracy score of training set with only Sex and Age is 0.7867564534231201
```

```
[20]: y_pred_sa = lr_sa.predict(X_test)
      pd.DataFrame({
      'PassengerId': test['PassengerId'],
      'Survived': y_pred_sa
      }).to_csv('Submit_Sex_Age.csv', index=False)
```

## OT3

```
[21]: print(X)

[[ 3.  0. 22.  0.]
 [ 1.  1. 38.  1.]
 [ 3.  1. 26.  0.]
 ...
 [ 3.  1. 28.  0.]
 [ 1.  0. 26.  1.]
 [ 3.  0. 32.  2.]]
```

normalized Age

```
[22]: mx_age, mn_age = X[:, 2].max(), X[:, 2].min()

      def normalize_age(x, mx_age, mn_age):
      return (x - mn_age) / (mx_age - mn_age)

      normalize_age_vectorized = np.vectorize(lambda x : normalize_age(x, mx_age,
      ↪mn_age))
      X[:, 2] = normalize_age_vectorized(X[:, 2])
      print(X)

      [[3.          0.          0.27117366  0.          ]
       [1.          1.          0.47222292  1.          ]
       [3.          1.          0.32143755  0.          ]
       ...
       [3.          1.          0.34656949  0.          ]
       [1.          0.          0.32143755  1.          ]
       [3.          0.          0.39683338  2.          ]]
```

```
[23]: class LinearRegressionGradient:
      def __init__(self, lr=0.001, random_state=42, epochs=200_000):
      self.lr = lr
      self.random_state = random_state
      self.epochs = epochs
      self.params = None
```

```

def fit(self, X: npt.ArrayLike, y: npt.ArrayLike):
    np.random.seed(self.random_state)
    X = np.hstack((np.ones(X.shape[0]).reshape(-1, 1), X)) # add bias
    self.params = np.random.randn(X.shape[1])

    for _ in range(self.epochs):
        y_pred = X @ self.params
        diff = y - y_pred
        loss = X.T @ diff

        self.params += self.lr / X.shape[0] * loss

    return self

def predict(self, X: npt.ArrayLike):
    X = np.hstack((np.ones(X.shape[0]).reshape(-1, 1), X)) # add bias
    return X @ self.params

```

```
[24]: params_gradient = LinearRegressionGradient(random_state=0).fit(X, y).params
      params_gradient
```

```
[24]: array([ 0.74253777, -0.18302629,  0.4945769 , -0.35394677,  0.04905888])
```

## OT4

```
[25]: class LinearRegressionInversion:
      def __init__(self):
          self.params = None

      def fit(self, X: npt.ArrayLike, y: npt.ArrayLike):
          X = np.hstack((np.ones(X.shape[0]).reshape(-1, 1), X)) # add bias

          self.params = np.linalg.inv(X.T @ X) @ X.T @ y
          return self

      def predict(self, X: npt.ArrayLike):
          X = np.hstack((np.ones(X.shape[0]).reshape(-1, 1), X)) # add bias

          return X @ self.params

```

```
[26]: params_matrix_inversion = LinearRegressionInversion().fit(X, y).params
      params_matrix_inversion
```

```
[26]: array([ 0.77442159, -0.18843944,  0.49086711, -0.40222591,  0.04911346])
```

Compute MSE

```
[27]: np.power(params_gradient - params_matrix_inversion, 2).sum()
```

```
[27]: 0.003390521142094702
```